

# Modern CPU Performance Profiling

**Hagen Paul Pfeifer**


ProtocolLabs

Hofmannstraße 19

81379 München, Germany

<http://www.protocollabs.com>


[hagen.pfeifer@protocollabs.com](mailto:hagen.pfeifer@protocollabs.com)



# Introduction

*„The purpose of computing is insight,  
not numbers.“*


— Richard W. Hamming




# Introduction

- No tool can be better than the user – this is especially true for performance analysis!
- Depending on the level of optimization you need a fairly good understanding of
  - How the compiler work
  - Processor specific instructions
  - How instructions are scheduled and the costs of instructions
  - About memory hierarchy: CPU caches, DRAM and system bus
  - IO Subsystem knowledge: network interface adapter, harddisk, ...
  - Kernel knowledge (including system calls, memory management, scheduling, IO system, ...)

# Introduction



# Introduction



# Table of Contents

- ① A Brief History of Time
- ② The TAO of Performance Analysis
- ③ Performance Monitoring Unit Evolution
- ④ Linux Perf
- ⑤ Examples

# A Brief History of Time



## Evolution of Latency and Bandwidth

	VAX 11 (1977)	Sandy Bridge	Improvement
Clock Speed	5 MHz <sup>1</sup>	3000 MHz	500x
Memory Size	2 MB	8000 MB	4000x
Memory Bandwidth	13 MB/s	5000 MB/s	385x
Memory Latency	225 ns	70 ns	3x

---

<sup>1</sup> 200 ns cycle time



# Latency and Bandwidth



# Latency and Bandwidth

- Bandwidth
  - How wide is your pipeline



# Latency and Bandwidth

- Bandwidth
  - How wide is your pipeline
- Latency
  - How long is your pipeline



# Latency and Bandwidth

- Bandwidth
  - How wide is your pipeline
- Latency
  - How long is your pipeline



You can solve bandwidth problems – but you can't solve latency problems!

# How can we cope with latency?

## How can we cope with latency?



## How can we cope with latency?

- SIMD - Single Instruction Multiple Data
  - Intel's MMX, SSE, SSE2, SSE3, SSSE3, SSE4.1, SSE4.2, ...
  - AMD's 3DNow
  - IBM's AltiVec
  - ARM's NEON
- Graphics Processing Units (GPU)

# How can we cope with latency?





## How can we cope with latency?

- Parallelize - process concurrently where ever possible:
  - Pipelining
  - Instruction Decoding
  - Hyperthreading
  - SMP/CMP Systems
- Cache<sup>2</sup> as much as possible
- Increase memory hierarchy
  - Uops cache
  - L1 cache
  - L2 cache
  - L3 cache
  - HD flash memory cache
  - HD

---

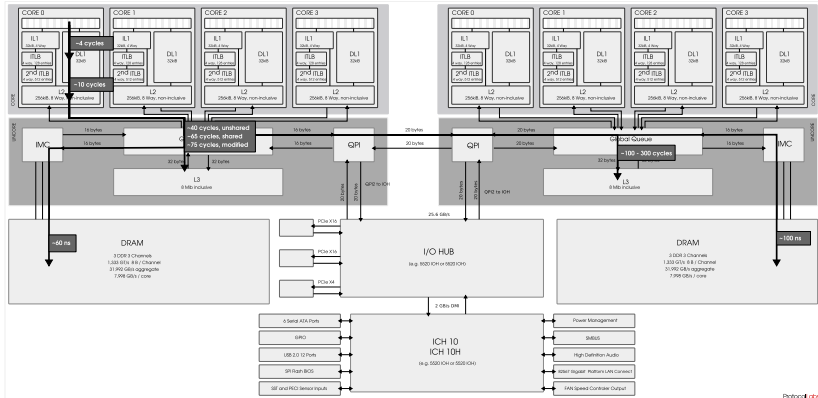
<sup>2</sup>SRAM is expensive; currently 95% of a core is already for the cache; in the end: you cannot increase SRAM to  $\infty$

## Bottom Line

All hardware vendor attempts have only one goal: *hide memory latency*



# Memory Hierarchy Today – Intel XEON 5000 Sequence



ProtocolLabs

# The Tao of Performance Analysis




# The Tao of Performance Analysis

- Identify hot spots
- Focus on hot spots, ignore the rest
- Reduce analyze complexity
  - Disable uninvolved services
  - Disable Hyperthreading, disable Turbo Boost
- Make analyze comparable: use same test-case
- Use realistic test-data

## Determine Efficiency of Hotspots

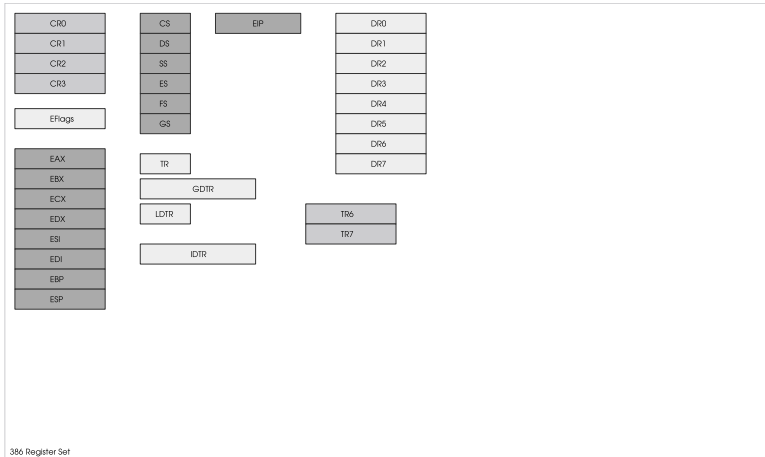
- Execution stalls
  - `UOPS_EXECUTED.CORE_STALL_CYCLES`,  
`UOPS_EXECUTED.CORE_ACTIVE_CYCLES`
  - Frontend stalls
  - Backend stalls
- Cycles per Instruction
  - `CPU_CLK_UNHALTED.CORE / INST_RETIRED.ANY`
  - Ideal CPI: 0.25, 1: OK, > 2: somehow high
- Examine generated instructions
  - Yes, instruction and architecture knowledge is required

# The Call for Hardware Performance Monitoring Support

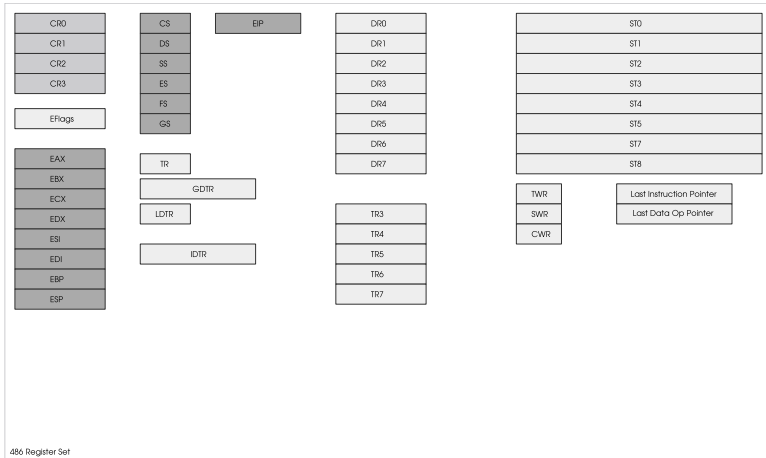




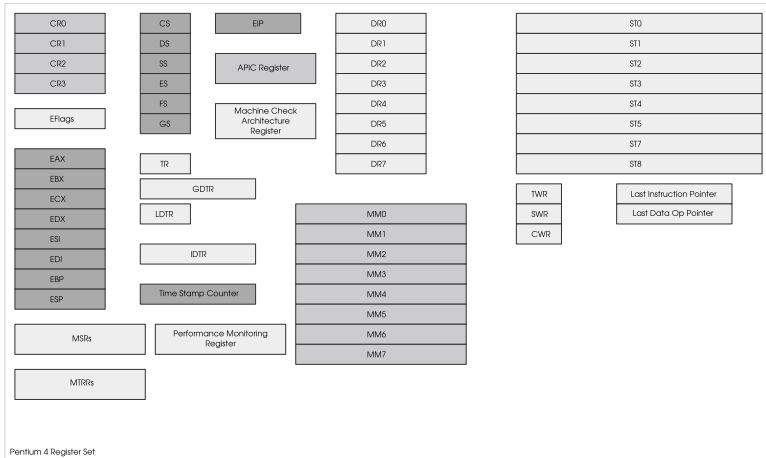
# Register Set Evolution



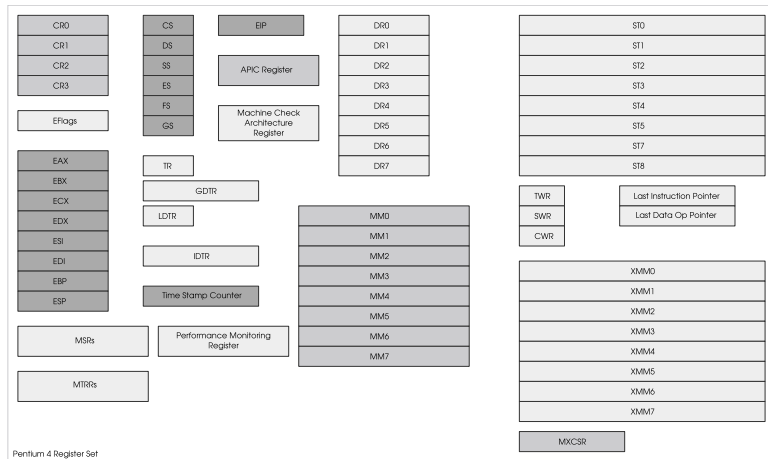
# Register Set Evolution



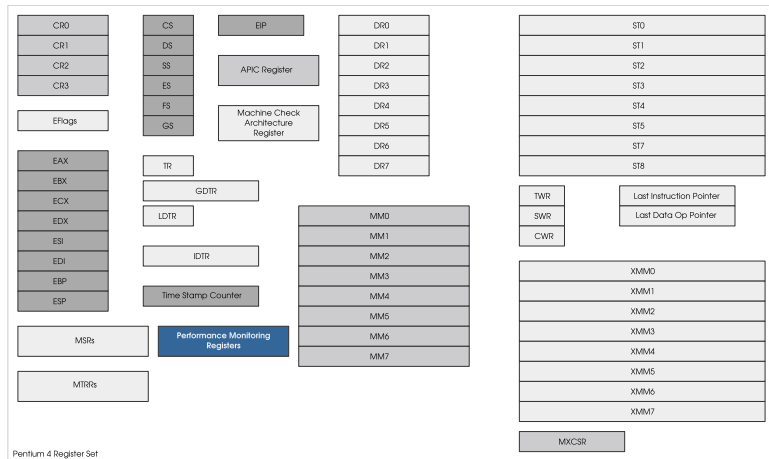
# Register Set Evolution



# Register Set Evolution



# Register Set Evolution



# Architectural Complexity

- With increasing architectural complexity it is almost impossible to know what happens „inside“.
- Modern CPU provides therefore insights and mechanism to get intrinsic data<sup>3</sup>

---

<sup>3</sup>or: current hardware trends makes monitoring critical

## Performance Monitoring Units – Overview

- One PMU in each (logical) Core
  - CPU\_CLK\_UNHALTED
  - INST\_RETIRE
  - LLC\_MISSES
  - LLC\_REFS
  - BR\_INST\_RETIRE
  - BR\_MISS\_PRED\_RETIRE
  - ...

## Performance Monitoring Units – Overview

- One PMU in each (logical) Core
  - CPU\_CLK\_UNHALTED
  - INST\_RETIRE
  - LLC\_MISSES
  - LLC\_REFS
  - BR\_INST\_RETIRE
  - BR\_MISS\_PRED\_RETIRE
  - ...
- One PMU in each Uncore
  - UNC\_DRAM\_PAGE\_MISS
  - UNC\_DRAM\_READ\_CAS
  - UNC\_GQ\_ALLOC
  - UNC\_GQ\_CYCLES\_FULL



## Performance Monitoring Units – Overview

- One PMU in each (logical) Core
  - CPU\_CLK\_UNHALTED
  - INST\_RETIRED
  - LLC\_MISSES
  - LLC\_REFS
  - BR\_INST\_RETIRED
  - BR\_MISS\_PRED\_RETIRED
  - ...
- One PMU in each Uncore
  - UNC\_DRAM\_PAGE\_MISS
  - UNC\_DRAM\_READ\_CAS
  - UNC\_GQ\_ALLOC
  - UNC\_GQ\_CYCLES\_FULL
- PMUs are highly specific to a CPU!

## Core PMU

- 4 fully programmable counters per CPU Core
  - Each count a certain event
  - Each counter can be used simultaneously
- 1 fixed counter
  - Clock Counter
  - Reference Clock Counter
  - Instruction Counter

# UNCore PMU

- 8 fully programmable counters per CPU Core
  - Each count a certain event
  - Each counter can be used simultaneously
- 1 fixed counter

# PMU Modes

- Event based Counting (EBC)
  - No interrupts
  - Very small overhead
- Event based Sampling (EBS)
  - Interrupts the CPU after a certain number of events<sup>4</sup>
  - Collect execution events
  - Statistical method
- Precise Event based Sampling (PEBS)
  - Not all events
  - Processor save (exact) context (instruction pointer)

---

<sup>4</sup> Sample After Value – SAV

# Linux Perf



## Perf Overview

- Ingo Molnar and Thomas Gleixner
- Milestone 8647093: from Documentation to new tool directory (2009)
- Userspace tool plus Kernel Subsystem<sup>5</sup>
- git like subcommands
- Per thread/per workload/per CPU/system wide
- No daemon mode
- Today: Arnaldo, Ingo, Frederic, Masami and many others

---

<sup>5</sup>maybe the term subsystem is to big, compared to real Linux subsystems

## Perf Modules

```
$ perf list
```

The most commonly used perf commands are:

<b>annotate</b>	Read perf.data (created by perf record) and display annotated code
<b>archive</b>	Create archive with object files with build-ids found in perf.data file
<b>bench</b>	General framework for benchmark suites
<b>buildid-cache</b>	Manage build-id cache.
<b>buildid-list</b>	List the buildids in a perf.data file
<b>diff</b>	Read two perf.data files and display the differential profile
<b>evlist</b>	List the event names in a perf.data file
<b>inject</b>	Filter to augment the events stream with additional information
<b>kmem</b>	Tool to trace/measure kernel memory(slab) properties
<b>kvm</b>	Tool to trace/measure kvm guest os
<b>list</b>	List all symbolic event types
<b>lock</b>	Analyze lock events
<b>probe</b>	Define new dynamic tracepoints
<b>record</b>	Run a command and record its profile into perf.data
<b>report</b>	Read perf.data (created by perf record) and display the profile
<b>sched</b>	Tool to trace/measure scheduler properties (latencies)
<b>script</b>	Read perf.data (created by perf record) and display trace output
<b>stat</b>	Run a command and gather performance counter statistics
<b>test</b>	Runs sanity tests.
<b>timechart</b>	Tool to visualize total system behavior during a workload
<b>top</b>	System profiling tool.

## Life Demo

# Demo



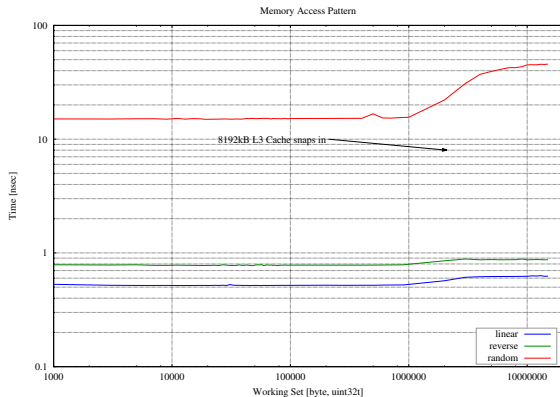
## Example 1

# Memory Access Pattern Matters

# Access Pattern

- Think about data structures
  - Array
  - List
  - Hashtable
  - Btree

# Access Pattern



## Access Pattern - Pseudo Random

139508,315690	task-clock	#	5,996 CPUs utilized	( +- 0,03% )	100,00%
12.673	context-switches	#	0,091 K/sec	( +- 0,24% )	100,00%
118	CPU-migrations	#	0,001 K/sec	( +- 13,19% )	100,00%
2.902	page-faults	#	0,021 K/sec	( +- 3,91% )	
78.854.265.280	cycles	#	0,565 GHz	( +- 1,10% )	40,58%
2.321.529.838	stalled-cycles-frontend	#	2,94% frontend cycles idle	( +- 4,28% )	38,57%
30.009.130.064	stalled-cycles-backend	#	38,06% backend cycles idle	( +- 0,12% )	42,99%
62.716.912.813	instructions	#	0,80 insns per cycle		
		#	0,48 stalled cycles per insn	( +- 4,37% )	40,51%
18.228.924.903	branches	#	130,666 M/sec	( +- 2,11% )	38,56%
65.516.732	branch-misses	#	0,36% of all branches	( +- 8,56% )	38,27%
31.569.908.790	L1-dcache-loads	#	226,294 M/sec	( +- 1,30% )	37,04%
18.172.049	L1-dcache-load-misses	#	0,06% of all L1-dcache hits	( +- 22,86% )	40,92%
50.183.885	LLC-loads	#	0,360 M/sec	( +- 14,13% )	40,83%
28.160.133	LLC-load-misses	#	56,11% of all LL-cache hits	( +- 7,80% )	41,73%
23,266494724	seconds time elapsed			( +- 0,03% )	

## Access Pattern - Linear

8042,299034	task-clock	#	5,998 CPUs utilized	( +- 0,05% )	100,00%
878	context-switches	#	0,109 K/sec	( +- 13,15% )	100,00%
24	CPU-migrations	#	0,003 K/sec	( +- 44,68% )	100,00%
2.646	page-faults	#	0,329 K/sec	( +- 1,76% )	
5.957.240.006	cycles	#	0,741 GHz	( +- 30,84% )	35,44%
1.419.666.336	stalled-cycles-frontend	#	23,83% frontend cycles idle	( +- 10,87% )	62,51%
1.096.067.572	stalled-cycles-backend	#	18,40% backend cycles idle	( +- 9,10% )	75,54%
3.087.516.698	instructions	#	0,52 insns per cycle		
		#	0,46 stalled cycles per insn	( +- 11,80% )	77,89%
1.567.151.858	branches	#	194,864 M/sec	( +- 17,85% )	59,37%
3.965.022	branch-misses	#	0,25% of all branches	( +- 30,58% )	31,14%
3.060.755.224	L1-dcache-loads	#	380,582 M/sec	( +- 22,17% )	19,07%
5.074.888	L1-dcache-load-misses	#	0,17% of all L1-dcache hits	( +- 46,08% )	9,92%
12.268.349	LLC-loads	#	1,525 M/sec	( +- 44,56% )	10,60%
2.566.738	LLC-load-misses	#	20,92% of all LL-cache hits	( +- 47,68% )	18,55%
1,340719886	seconds		time elapsed	( +- 0,05% )	

## Example 2

# Shared Resources

## Example 2

**Input:** global variable  $i$   
pthread\_barrier\_wait(&barrier);  
**for**  $i \leftarrow 0$  **to**  $2^{30}$  **do**  
     $i++$ ;  
**end**

**Input:** global variable  $i$   
pthread\_barrier\_wait(&barrier);  
**for**  $i \leftarrow 0$  **to**  $2^{30}$  **do**  
     $i++$ ;  
**end**

## Shared Resources


```
$ perf stat -d -r 10 ./a.out
197170,885666 task-clock                #    5,318 CPUs utilized          ( +-  3,90% )
      12.611 context-switches          #    0,064 K/sec                  ( +- 20,75% )
        93 CPU-migrations              #    0,000 K/sec                  ( +- 22,33% )
       207 page-faults                 #    0,001 K/sec                  ( +-  0,18% )
642.275.775.062 cycles                  #    3,257 GHz                    ( +-  3,90% ) 40,00%
  14.223.466.981 stalled-cycles-frontend #    2,21% frontend cycles idle   ( +-  2,44% ) 40,01%
608.715.881.283 stalled-cycles-backend  #   94,77% backend cycles idle    ( +-  3,98% ) 40,00%
   4.020.392.445 instructions           #    0,01 insns per cycle         ( +-  1,02% ) 40,01%
                                # 151,41 stalled cycles per insn  ( +-  1,25% ) 40,02%
      692.627.784 branches              #    3,513 M/sec                  ( +- 14,30% ) 40,01%
     5.289.831 branch-misses            #    0,76% of all branches        ( +-  2,59% ) 40,01%
14.730.027.536 L1-dcache-loads          #   74,707 M/sec                  ( +-  2,72% ) 40,01%
   392.392.754 L1-dcache-load-misses    #    2,66% of all L1-dcache hits  ( +-  2,64% ) 40,01%
   1.088.906.661 LLC-loads              #    5,523 M/sec                  ( +-  2,76% ) 40,00%
   1.015.331.233 LLC-load-misses        #   93,24% of all LL-cache hits   ( +-  2,64% )

37,072694341 seconds time elapsed      ( +-  2,64% )
```



# Pipelining

- Modern processors employ a technique called pipelining to increase instruction throughput
- Various pieces of hardware perform various operation at the same time (parallel)
  - Part A of the pipeline is executing instruction FOO
  - Part B fetch instruction Bar
  - Part C decode instruction X
  - Part D commit results from instruction FOO



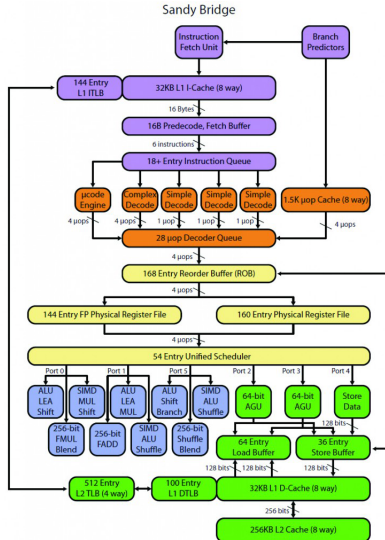
# Frontend

- Responsible for providing a stream of work for the back-end
- Front-end operate "in-order"
- Working on instructions (but generate uops for backend, called decoding)
- Sandy Bridge Front-end is capable of delivering up to 4 instructions/cycle (4 decoders) to the back.

# Backend

- Front-end operate "out-of-order" (often)
- Working on micro-operations (uops)

# Sandy Bridge Microarchitecture



## Find the hot spot

```
perf record -a -g -e stalled-cycles-backend ./a.out  
perf report
```

# Find the hot spot

```
Samples: 814K of event 'stalled-cycles-backend', Event count (approx.): 627665590172
+ 18,35% false-sharing false-sharing [.] t2
+ 18,32% false-sharing false-sharing [.] t4
+ 18,24% false-sharing false-sharing [.] t3
+ 17,80% false-sharing false-sharing [.] t1
+ 13,21% false-sharing false-sharing [.] t6
+ 12,69% false-sharing false-sharing [.] t5
+ 1,23% swapper [kernel.kallsyms] [k] native_safe_halt
+ 0,02% Xorg [kernel.kallsyms] [k] delay_tsc
+ 0,02% Xorg [kernel.kallsyms] [k] in_lock_functions
+ 0,01% Xorg [kernel.kallsyms] [k] native_read_tsc
```

# Find the hot spot

t2		
		sub \$0x18,%rsp
		mov \$0x601c20,%edi
	→	callq pthread_barrier_wait@plt
		movq \$0x0,0x8(%rsp)
	↓	jmp 2e
57,13	19:	mov num2,%eax
0,01		add \$0x1,%eax
42,85		mov %eax,num2
		addq \$0x1,0x8(%rsp)
0,00	2e:	cmpq \$0x5f5e0ff,0x8(%rsp)
0,01	↑	jbe 19
		mov num2,%eax
		add \$0x18,%rsp
	←	retq

## False Sharing

- False sharing is when different threads access non-overlapping areas of a cachline
- Here: one cachline is shared between 6 cores!
- Causing lines to be renewed regularly, if any thread writes to it
- BTW: Linux Kernel have PER\_CPU  
(`DEFINE_PER_CPU(type, name), ...`)



# Try to fix this

```
$ diff -Nuar false-sharing.c.old false-sharing.c
--- false-sharing.c.old 2012-05-23 22:07:52.000000000 +0000
+++ false-sharing.c      2012-05-23 22:08:16.000000000 +0000
@@ -33,12 +33,12 @@

static pthread_barrier_t barrier;

-static unsigned num1;
-static unsigned num2;
-static unsigned num3;
-static unsigned num4;
-static unsigned num5;
-static unsigned num6;
+static unsigned num1 ____cacheline_aligned;
+static unsigned num2 ____cacheline_aligned;
+static unsigned num3 ____cacheline_aligned;
+static unsigned num4 ____cacheline_aligned;
+static unsigned num5 ____cacheline_aligned;
+static unsigned num6 ____cacheline_aligned;

$ make
$ perf stat -a -d -r 10 ./false-sharing
1257,529627 task-clock # 6,011 CPUs utilized ( +- 0,82% ) 100,00%
422 context-switches # 0,335 K/sec ( +- 1,42% ) 100,00%
22 CPU-migrations # 0,017 K/sec ( +- 4,24% ) 100,00%
212 page-faults # 0,169 K/sec ( +- 0,18% )
3.701.493.777 cycles # 2,943 GHz ( +- 0,98% ) 40,82%
190.292.288 stalled-cycles-frontend # 5,14% frontend cycles idle ( +- 6,17% ) 41,67%
2.207.591.340 stalled-cycles-backend # 59,64% backend cycles idle ( +- 0,72% ) 41,52%
3.623.979.984 instructions # 0,98 insns per cycle
# 0,61 stalled cycles per insn ( +- 0,49% ) 40,91%
630.872.872 branches # 501,676 M/sec ( +- 0,90% ) 39,90%
270.216 branch-misses # 0,04% of all branches ( +- 7,45% ) 38,93%
2.547.084.258 L1-dcache-loads # 2025,467 M/sec ( +- 1,03% ) 38,56%
337.080 L1-dcache-load-misses # 0,01% of all L1-dcache hits ( +- 13,59% ) 38,85%
509.484 LLC-loads # 0,405 M/sec ( +- 10,60% ) 39,17%
120.115 LLC-load-misses # 23,58% of all LL-cache hits ( +- 9,60% ) 39,82%

0,209195268 seconds time elapsed ( +- 0,82% )
```

# SNAFU

```

t1
    sub    $0x18,%rsp
    mov    $0x601c40,%edi
    → callq pthread_barrier_wait@plt
    movq   $0x0,0x8(%rsp)
    ↓ jmp   2e
10,56  19:  mov    num1,%eax
27,35   add    $0x1,%eax
        mov    %eax,num1
16,54   addq   $0x1,0x8(%rsp)
35,50  2e:  cmpq   $0x5f5e0ff,0x8(%rsp)
10,05   ↑ jbe   19
        mov    num1,%eax
        add    $0x18,%rsp
        ← retq

```

## Suggested Readings

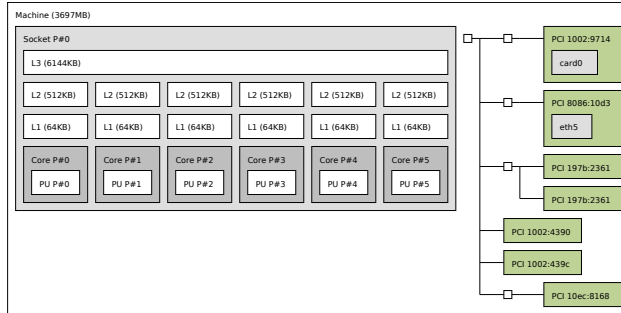
- Books
  - *Inside the Machine: An Illustrated Introduction to Microprocessors and Computer Architecture* from Jon Stokes
  - *Multi-Core Programmierung* Intel Press
- Online Articles:
  - Intel Guide for Developing Multithreaded Applications  
<http://software.intel.com/en-us/articles/intel-guide-for-developing-multithreaded-applications/>
  - The Architecture of the Nehalem Processor and Nehalem-EP SMP Platforms, Michael E. Thomadakis, Supercomputing Facility, Texas A&M University
  - BIOS and Kernel Developer's Guide for AMD NPT Family 0Fh Processors
  - Intel XEON Processor 7500 Series Uncore Programming Guide
  - Intel 5520 Chipset and Intel 5500 Chipset
  - Intel Microarchitecture Codename Nehalem Performance Monitoring Unit Programming Guide  
<http://software.intel.com/file/30320>
  - Performance Analysis Guide for Intel Core i7 Processor and Intel Xeon 5500 processor

Thank you very much!

Questions?

# Modern Memory Architecture

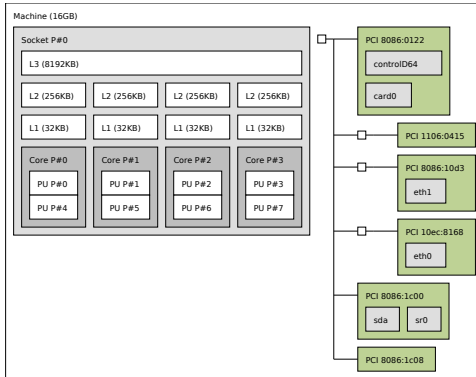
## AMD Phenom(tm) II X6 1055T Processor<sup>6</sup>



<sup>6</sup> generated via `lstopo --no-legend mm-arch-x2-1055.pdf`; Debian package: *hwloc*

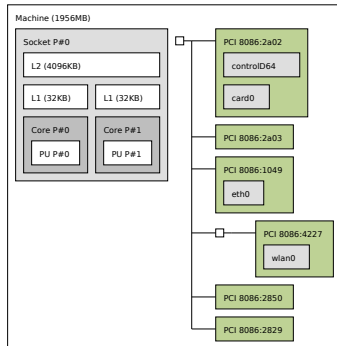
# Modern Memory Architecture

Intel(R) Core(TM) i7-2700K CPU @ 3.50GHz



# Modern Memory Architecture

Intel(R) Core(TM)2 Duo CPU T7500 @ 2.20GHz




# Intel Sandy Bridge Architecture

Target segment	Cores (Threads)	Processor Branding & Model	CPU Clock rate		Graphics Clock rate		L3 Cache	TDP	Release Date (Y-M-D)	Price (USD)	Motherboard		
			Normal	Turbo	Normal	Turbo					Socket	Interface	Memory
Extreme / High-End	6 (12)	Core i7 Extreme	3960X	3.3 GHz	3.9 GHz	N/A	15 MB	130 W	2011-11-14	\$999	LGA 2011	DMI 2.0 PCIe 2.0 <sup>[14]</sup>	Up to quad channel DDR3-1600 <sup>[15]</sup>
			2930K	3.2 GHz	3.8 GHz		12 MB		2012-02-13 <sup>[36]</sup>	\$583			
			3820	3.6 GHz			10 MB		2011-10-24	\$332			
Performance	4 (8)	Core i7	2700K	3.5 GHz	3.9 GHz	850 MHz 1350 MHz	8 MB	95 W	2011-01-09	\$294	LGA 1155	DMI 2.0 PCIe 2.0	Up to dual channel DDR3-1333
			2600K	3.4 GHz					2011-01-09	\$317			
			2600		3.8 GHz				2011-01-09	\$294			
			2600S	2.8 GHz					2011-01-09	\$306			
			2550K	3.4 GHz					2012-01-30	\$225			
			2500K						2012-01-30	\$216			
	4 (4)	Core i5	2500	3.3 GHz	3.7 GHz	850 MHz 1100 MHz	6 MB	95 W	2011-09-04	\$205			
			2500S	2.7 GHz					2011-09-04	\$216			
			2500T	2.3 GHz	3.3 GHz	650 MHz 1250 MHz			2011-05-22	\$205			
			2450P	3.2 GHz	3.5 GHz	N/A		95 W	2012-01-30	\$195			
			2400	3.1 GHz	3.4 GHz				2011-01-09	\$184			
			2405S	2.5 GHz	3.3 GHz	850 MHz 1100 MHz			2011-05-22	\$205			
			2400S					95 W	2011-01-09	\$195			
			2380P	3.1 GHz	3.4 GHz	N/A			2012-01-30				
			2320	3.0 GHz	3.3 GHz				2011-09-04	\$177			
			2310	2.9 GHz	3.2 GHz	850 MHz		95 W	2011-05-22				
			2300	2.8 GHz	3.1 GHz				2011-01-09				
			2390T	2.7 GHz	3.5 GHz				2011-02-20	\$195			
			2120T	2.6 GHz		650 MHz		35 W	2011-09-04	\$127	LGA 1155	DMI 2.0 PCIe 2.0	
			2100T	2.5 GHz					2011-02-20				
			2130	3.4 GHz						\$138			




# Intel Sandy Bridge Architecture

- Replace Nehalem microarchitecture
- Released first January 2011, developed beginning in 2005
- Major features
  - 32 kB data + 32 kB instruction L1 cache (3 clocks) and 256 kB L2 cache (8 clocks) per cor
  - 64-byte cache line size
  - Shared L3 cache includes the processor graphics
  - Decoded micro-operation cache and enlarged, optimized branch predictor
  - Two load/store operations per CPU cycle for each memory channel
- Roadmap:



# Cache Evolution



source: anandtech.com

Hagen Paul Pfeifer

May 20, 2012

53/58

## Use Latest Perf Version

- You can use Arnaldo's or Peter's branch directly (sometimes I do), but I suggest to use Ingo's development master branch! (tip:master)
- You are free to use perf/core from Ingo too.

```
git remote add tip git://git.kernel.org/pub/scm/linux/kernel/git/tip/tip.git
git remote update tip
git checkout -b perf/core tip/perf/core
or
git checkout -b tip/master tip/master
```

# Latency Values

- Mutex lock/unlock 100 ns
- Send 2K bytes over 1 Gbps network 20,000 ns
- Read 1 MB sequentially from memory 250,000 ns
- Disk seek 10,000,000 ns
- ...

## Perfmon2


- Alternative for perf
- Support sampling and counting
- System/Thread wide
- Perfmon provides a library, useable for self-monitoring!
- ...

## Perfmon2

- Alternative for perf
- Support sampling and counting
- System/Thread wide
- Perfmon provides a library, useable for self-monitoring!
- ...
- Great tool to understand PMUs in detail!

<http://perfman2.sf.net>

# Data Tiling - Cache Blocking



## Thread - CPU Detail

- `watch -n ,1 ps -aLo pcpu,cpuid,pid,args`